

# Caffe V1 Examples

- LeNet and YOLO -

Aug. 2019

Ando Ki, Ph.D.

[adki@future-ds.com](mailto:adki@future-ds.com)

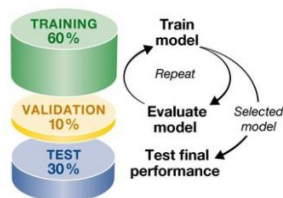
## LeNet Example

- LeNet-5 for MNIST
- LeNet-5 for MNIST: layer
- LeNet-5 for MNIST: all together
- LeNet-5 for MNIST: running
- LeNet-5 for MNIST: solver
- LeNet-5 for MNIST: net
- Running LeNet with Caffe
- Run inference with sample image
- Deploy prototxt
- Caffe Python interface for LeNet

# LeNet-5 for MNIST

## ■ MNIST dataset

- ▶ Modified National Institute of Standards and Technology
- ▶ Handwritten digits database
  - ⇒ 10 classes: 0, 1, ..., 9
  - ⇒ training set: 60,000 training image
  - ⇒ test set: 10,000 testing image

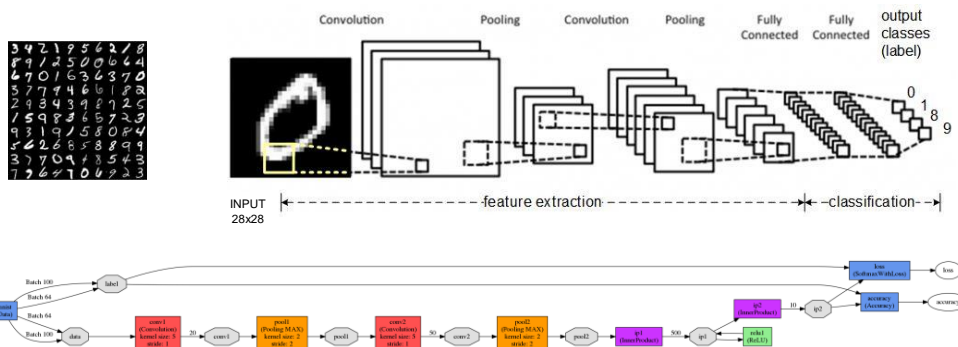


3

# LeNet-5 for MNIST

## ■ LeNet is one of the popular convolutional networks, and works well on digit classification tasks.

- ⇒ 784 (28x28) inputs of black and white → converted to floating number 0.0 ~ 1.0
- ⇒ 10 outputs representing digit 0 to 9




4

## LeNet-5 for MNIST: layer


### 1<sup>st</sup> convolution layer

- ▶ Input: 28x28 pixels
- ▶ Convolution filter: 32 kernels with 5x5
- ▶ Convolution: stride 1
  - ☞ It generates the same number of elements
- ▶ Results in: 32 features of 28x28



### 1<sup>st</sup> pooling layer (sub-sampling)

- ▶ Input: 32 features with 28x28
- ▶ Max pooling filter: 5x5 (2x2 ?)
- ▶ Convolution: stride 2
  - ☞ It generates ½ number of elements
- ▶ Results in: 32 features of 14x14




5

## LeNet-5 for MNIST: layer

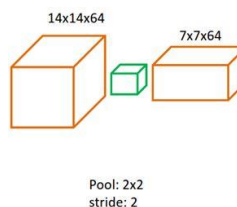
### 2<sup>nd</sup> convolution

- ▶ Input: 32 features with 14x14 pixels
  - ☞ 32 kernels are used at the previous stage
- ▶ Convolution filter: 64 kernels with 5x5
- ▶ Convolution: stride 1
  - ☞ It generates the same number of elements
- ▶ Results in: 64 features of 14x14



### 2<sup>nd</sup> pooling

- ▶ Input: 64 features with 14x14
- ▶ Max pooling filter: 2x2
- ▶ Convolution: stride 2
  - ☞ It generates ½ number of elements
- ▶ Results in: 64 features of 7x7



6



## LeNet-5 for MNIST: running

### Steps (in details)

- ▶ go to project directory
  - `$ cd work/codes/caffe_v1-projects/mnist.LeNet`
- ▶ get dataset:
  - `$ ./scripts/get_mnist.sh data`
  - (ungzip all in 'data' directory)
- ▶ convert the dataset to Caffe data format
  - `$ ./scripts/create_mnist.sh ${CAFFE_HOME} data`
- ▶ training
  - `$ ./scripts/train_lenet.sh ${CAFFE_HOME} prototxt/lenet_solver.prototxt`
- ▶ running LeNet model with 'mnist\_test\_lmdb'
  - `$ ./scripts/test_lenet.sh`

### Step in simple

- ▶ go to project directory
  - `$ cd work/codes/caffe_v1-projects/mnist.LeNet`
- ▶ Run make
  - `$ make cleanall`
  - `$ make lmdb`
  - `$ make train`
  - `$ make test`

```

[osboxes@osboxes] make run
[export GLIBC_TUNING=ld-single]
[scripts/test_lenet.sh /home/osboxes/work/caffe/
prototxt/lenet_train_test.prototxt]
[0905 01:51:14.047650 3950 caffe.cpp:275] Use CPU.
[0905 01:51:14.049351 3950 net.cpp:290] The NetState phase (1) differed from th
e phase (0) specified by a rule in layer mnist
[0905 01:51:14.049351 3950 net.cpp:53] Initializing net from parameters:
name: "lenet"
state {
  solver_mode: "test"
[0905 01:51:17.037341 3950 caffe.cpp:304] Batch 47, accuracy = 0.957
[0905 01:51:17.037655 3950 caffe.cpp:304] Batch 47, loss = 0.0546983
[0905 01:51:17.037647 3950 caffe.cpp:304] Batch 48, accuracy = 0.956
[0905 01:51:17.037147 3950 caffe.cpp:304] Batch 48, loss = 0.140428
[0905 01:51:17.756855 3950 caffe.cpp:304] Batch 49, accuracy = 1
[0905 01:51:17.756772 3950 caffe.cpp:304] Batch 49, loss = 0.032333
[0905 01:51:17.756825 3950 caffe.cpp:309] Loss: 0.0872644
[0905 01:51:17.756883 3950 caffe.cpp:321] accuracy = 0.9706
[0905 01:51:17.756943 3950 caffe.cpp:321] loss = 0.0872644 (* 1 = 0.0872644 los
s)
[osboxes@osboxes]

```

9

## LeNet-5 for MNIST: solver

- ▶ net: network mode
- ▶ test\_iter: iterations to test
- ▶ test\_interval: interval between test
- ▶ base\_lr: Learning Rate initial value
- ▶ display: iterations to show progress
- ▶ max\_iter: max iterations for training.
- ▶ snapshot: iterations to store snapshot.
- ▶ solver\_mode: CPU or GPU

```

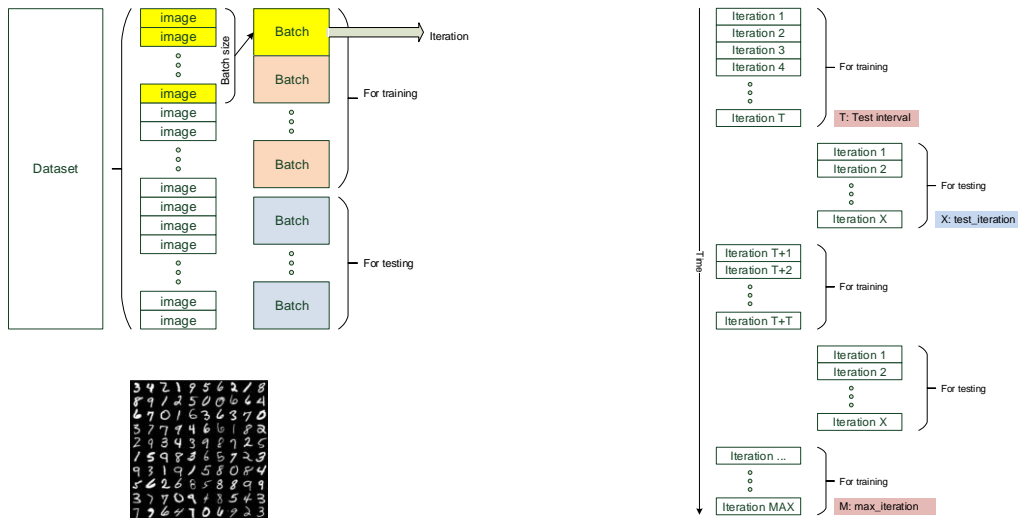
# MNIST lenet_solver.prototxt
net: "lenet_train_test.prototxt"
test_iter: 100
test_interval: 500
base_lr: 0.01
momentum: 0.9
weight_decay: 0.0005
lr_policy: "inv"
gamma: 0.0001
power: 0.75
display: 100
max_iter: 10000
snapshot: 5000
snapshot_prefix: "snapshots"
solver_mode: CPU

```

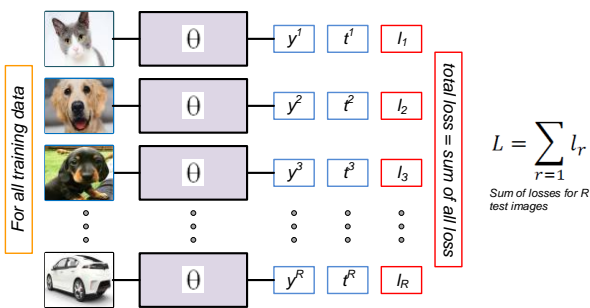


<https://github.com/BVLC/caffe/wiki/Solver-Prototxt>

10

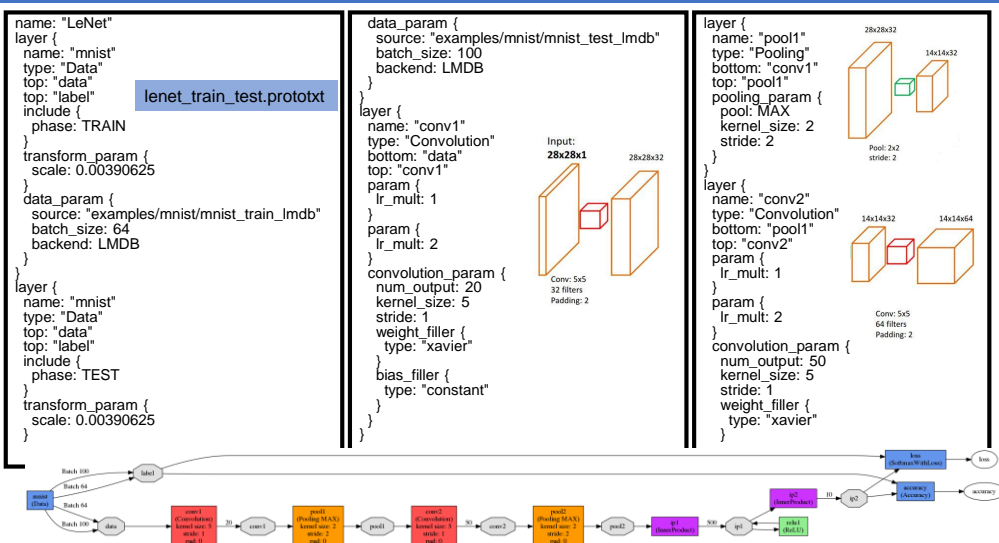
# LeNet-5 for MNIST: solver



# LeNet-5 for MNIST: solver

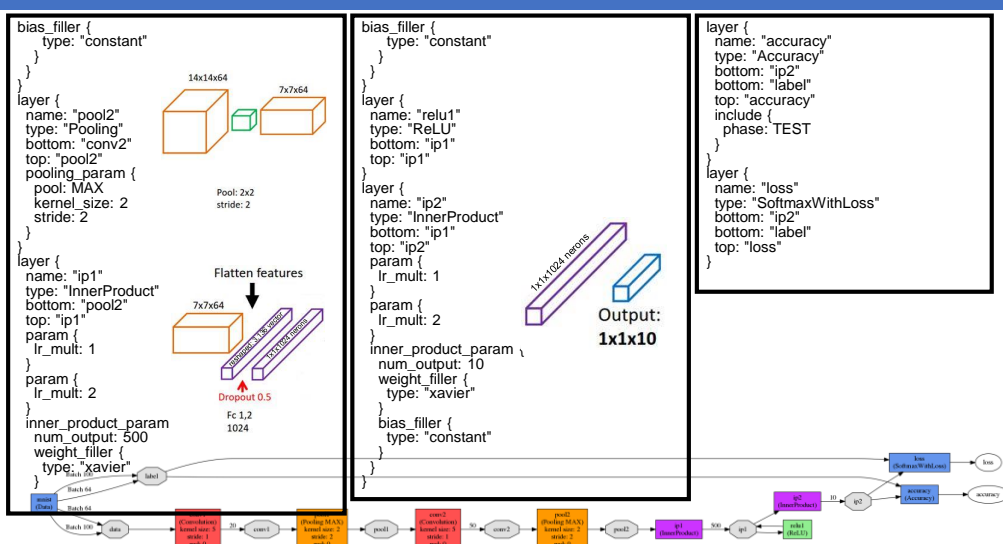


# LeNet-5 for MNIST: net



13

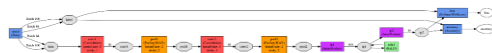
# LeNet-5 for MNIST: net



14

## Running LeNet with Caffe

- This example is about LeNet
- Make sure 'work/caffe' is ready
  - ▶ see the pervious slides
  - ▶ Step 1: go to your project directory
    - ☞ [user@host] cd \$(PROJECT)/codes.caffe/mnist.LeNet
  - ▶ Step 2: check network
    - ☞ [user@host] make draw
    - ☞ [user@host] fim lenet\_train\_test.png
  - ▶ Step 3: make data (convert data)
    - ☞ [user@host] make lmdb
  - ▶ Step 4: run train (it takes time)
    - ☞ [user@host] make train
  - ▶ Step 5: run loss graph
    - ☞ [user@host] make plot
  - ▶ Step 6: run test
    - ☞ [user@host] make test
  - ▶ Step 7: run deployment (inference)
    - ☞ [user@host] make deploy



```
Terminal
[osboxes@osboxes] make run
(export GLOG_minloglevel=0)
scripts/test_lenet.sh /home/osboxes/work/caffe/
prototxt/lenet_train_test.prototxt\
snapshots_iter_1000.caffemodel)
10905 01:51:14.647650 3950 caffe.cpp:275] Use CPU.
10905 01:51:14.649165 3950 net.cpp:296] The NetState phase (1) differed from th
e phase (0) specified by a rule in layer mnist
10905 01:51:14.649351 3950 net.cpp:53] Initializing net from parameters:
name: "LeNet"
state {
  TRAIN - TEST
10905 01:51:17.037341 3950 caffe.cpp:304] Batch 47, accuracy = 0.97
10905 01:51:17.637655 3950 caffe.cpp:304] Batch 47, loss = 0.0540983
10905 01:51:17.697047 3950 caffe.cpp:304] Batch 48, accuracy = 0.96
10905 01:51:17.697147 3950 caffe.cpp:304] Batch 48, loss = 0.140428
10905 01:51:17.756655 3950 caffe.cpp:304] Batch 49, accuracy = 1
10905 01:51:17.756772 3950 caffe.cpp:304] Batch 49, loss = 0.032333
10905 01:51:17.756825 3950 caffe.cpp:309] Loss: 0.0872644
10905 01:51:17.756833 3950 caffe.cpp:321] accuracy = 0.9766
10905 01:51:17.756943 3950 caffe.cpp:321] loss = 0.0872644 (* 1 = 0.0872644 los
s)
[osboxes@osboxes]
```


use 'display' for Ubuntu, 'fim' for Raspbian' to display image.

15

## Run inference with sample image

- Go to 'mnist.LeNet' directory
  - ▶ \$ cd ../codes/caffe-v1-projects/mnist.LeNet
- Run make
  - ▶ \$ make deploy
- Note that LeNet uses inverted image, i.e., background should be black.

```
Terminal
[osboxes@osboxes] make
GLOG_minloglevel=2 python mnist_test.py samples/4.png
[ 2.50871176e-06 1.81367841e-06 1.57160230e-05 2.03305008e-05
 9.89796937e-01 9.22584604e-06 5.09644167e-07 4.18145180e-04
 5.32380818e-06 9.72963311e-03]
4
[osboxes@osboxes]
```



16



## Deploy prototxt (1/2)

- Refer to 'lenet\_deploy.prototxt' under 'prototxt' directory.

<pre>##### Remove the data layer #layer { #  name: "mnist" #  type: "Data" #  top: "data" #  top: "label" #  include { #    phase: TRAIN #  } #  transform_param { #    scale: 0.00390625 #  } #  data_param { #    source: "data/mnist_train_lmdb" #    batch_size: 64 #    backend: LMDB #  } #}  ##### Remove the label layer #layer { #  name: "mnist" #  type: "Data" #  top: "data" #  top: "label" #  include { #    phase: TEST #  } #  transform_param { #    scale: 0.00390625 #  } #  data_param { #    source: "data/mnist_test_lmdb" #    batch_size: 100 #    backend: LMDB #  } #}</pre>	<pre>##### Add a new layer to accept data without label ### It define the name and shapes of the input blobs. # shape: { dim: 1 # batchsize (how many images/samples are fed through the #               network in parallel) #         dim: 28 # height of data, i.e., pixels (MNIST data is 28x28) #         dim: 28 # width of data, i.e., pixels (MNIST data is 28x28) #       } # layer { #   name: "data" #   type: "input" #   top: "data" #   input_param { #     shape: { dim: 1 #              dim: 28 #              dim: 28 #            } #   } # }  layer {   name: "conv1"   type: "Convolution"   bottom: "data"   top: "conv1"   param {     lr_mult: 1   }   param {     lr_mult: 2   } }  ##### other hidden layers remains</pre>
---	--

17

## Deploy prototxt (2/2)

- Refer to 'lenet\_deploy.prototxt' under 'prototxt' directory.

```
...
...
}
}

##### Remove the layers depending upon data labels
##### accordingly remove layer that uses 'data' as bottom
#layer {
#  name: "accuracy"
#  type: "Accuracy"
#  bottom: "ip2"
#  bottom: "label"
#  top: "accuracy"
#  include {
#    phase: TEST
#  }
#}

#layer {
#  name: "loss"
#  type: "SoftmaxWithLoss"
#  bottom: "ip2"
#  bottom: "label"
#  top: "loss"
#}

##### Add a new layer to the end of this network to produce a Softmax
output
layer {
  name: "loss"
  type: "Softmax"
  bottom: "ip2"
  top: "loss"
}
```

18

## Run inference with sample image (another way)

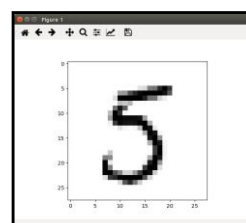
- Go to 'mnist.LeNet.python' directory

- ▶ \$ cd ../codes/caffe\_v1-project/mnist.LeNet.python

- Run make

- ▶ \$ make

```
Terminal
[osboxes@osboxes] make
GLOG_minloglevel=2 python mnist_test.py samples/4.png
[ 2.50871176e-06  1.81367841e-06  1.57160230e-05  2.03305008e-05
 9.89796937e-01  9.22584604e-06  5.09644167e-07  4.18145180e-04
 5.32380818e-06  9.72963311e-03]
4
[osboxes@osboxes]
```



19

## Caffe Python interface for LeNet

```
import os
os.environ['GLOG_minloglevel']='2'

import caffe
import numpy as np
import cv2
import sys
import Image

caffe_home = os.environ['CAFFE_ROOT'];
model = caffe_home + '/examples/mnist/lenet.prototxt';
weights = './mnist.LeNet/snapshots_iter_1000.caffemodel';
net = caffe.Net(model, weights, caffe.TEST);
caffe.set_mode_cpu()

img = cv2.imread(sys.argv[1],0)
if img.shape != [28,28]:
    img2 = cv2.resize(img,(28,28))
    img = img2.reshape(28,28,-1);
else:
    img = img.reshape(28,28,-1);

img = 1.0 - img/255.0

out = net.forward_all(data=np.asarray([img.transpose(2,0,1)]))

print out['prob'][0]
print out['prob'][0].argmax()
```

Using pycaffe

This may not need.

Revert image and normalize it to 0~1

Inference

Get the highest probability one

20

# YOLO

- Darknet: The framework for YOLO
- Darknet example for detector
- Network configuration file
- YOLO
- YOLO model
- YOLO network
- YOLO on Caffe

21

## Darknet: The framework for YOLO

- Darknet is an open source neural network framework written in C and CUDA (Compute Unified Device Architecture) supporting CPU (Central Processing Unit) and GPU (Graphical Processing Unit) computation.
  - ▶ <https://github.com/pjreddie/darknet>
  - ▶ <https://pjreddie.com/darknet/>
  - ▶ <https://groups.google.com/forum/#!forum/darknet>
  - ▶ Site: <https://pjreddie.com/darknet/>
  - ▶ GitHub : <https://github.com/pjreddie/darknet>

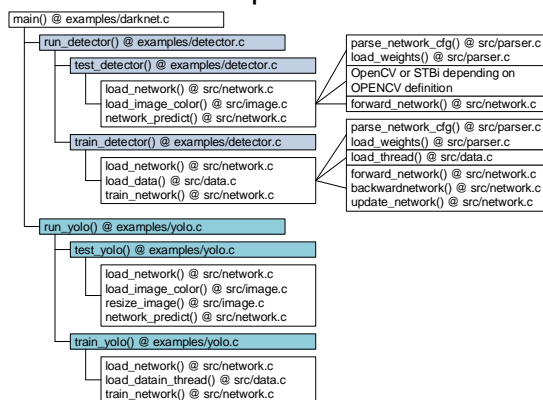


"Darknet: Open Source Neural Networks in C", Joseph Redmon, <http://pjreddie.com/darknet>, 2013-2016.

22

## Darknet example for detector

### ■ Have a look at 'darknet/examples/darknet.c' file



23

## Network configuration file

### ■ Nodes

- ▶ [shortcut]
- ▶ [crop]
- ▶ [cost]
- ▶ [detection]
- ▶ [region]
- ▶ [local]
- ▶ [conv] or [convolutional]
- ▶ [deconv] or [deconvolutional]
- ▶ [activation]
- ▶ [net] or [network]
- ▶ [crnn]
- ▶ [gru]
- ▶ [lstm]
- ▶ [rnn]
- ▶ [conn] or [connected]
- ▶ [max] or [maxpool]
- ▶ [reorg]

<https://github.com/cvjena/darknet/blob/master/cfg/yolo.cfg>

- ▶ [avg] or [avgpool]
- ▶ [dropout]
- ▶ [ln] or [normalization]
- ▶ [batchnorm]
- ▶ [soft] or [softmax]
- ▶ [route]

```

[net]
batch=1
subdivisions=1
width=416
height=416
channels=3

[convolutional]
batch_normalize=1
filters=16
size=3
stride=1
pad=1
activation=leaky

[maxpool]
size=2
stride=2

```

24

## Network configuration file

- [net] node
  - ▶ batch: how many images are in each batch to average the loss over?
  - ▶ subdivisions: into how many sub-batches shall each batch be divided to handle images in each sub-batch in parallel?
  - ▶ height, width: input size of the network
  - ▶ channels: number of components, e.g., color components
  - ▶ momentum: learning parameters
  - ▶ learning\_rate: base learning rate
  - ▶ policy: change learning rate after the corresponding steps
  - ▶ steps: need to have as many steps as scale
  - ▶ scales: re-scale the current learning rate by the corresponding factor once the number of steps is reached
  - ▶ max\_batches: max number of "iterations"
  - ▶ i\_snapshot\_iteration: snapshot the learned weights after every k "iterations"
- [convolutional] node
  - ▶ filters: number of filters, i.e., kernels
  - ▶ size: size of filter, e.g., 3 means 3x3 filter
  - ▶ stride: number of stride
  - ▶ pad: number of padding, e.g., 1
  - ▶ activation: specify activation function
- [maxpool] node
  - ▶ size: size of filter
  - ▶ stride: number of stride
- [connected] node
  - ▶ output: number of output of fully connected network
  - ▶ activation: activation function
- [detection] node or [region] node
  - ▶ classes: number of classes, e.g., 20 for pascal voc (l.classes)
  - ▶ coords: bounding boxes -> 4 parameters (l.n)
  - ▶ side: number of cell in x and y direction
  - ▶ num: number of predicted boxes per cell

25

## Network configuration file

- First node should be [net]

tiny-yolo.cfg									
layer	filters	size	input	output					
0 conv	16	3 x 3 / 1	416 x 416 x 3	-> 416 x 416 x 16	[net]	width=416			
1 max	2	2 x 2 / 2	416 x 416 x 16	-> 208 x 208 x 16		height=416			
2 conv	32	3 x 3 / 1	208 x 208 x 16	-> 208 x 208 x 32		channels=3			
3 max	2	2 x 2 / 2	208 x 208 x 32	-> 104 x 104 x 32	[convolutional]	batch_normalize=1			
4 conv	64	3 x 3 / 1	104 x 104 x 32	-> 104 x 104 x 64		filters=16			
5 max	2	2 x 2 / 2	104 x 104 x 64	-> 52 x 52 x 64		size=3			
6 conv	128	3 x 3 / 1	52 x 52 x 64	-> 52 x 52 x 128		stride=1			
7 max	2	2 x 2 / 2	52 x 52 x 128	-> 26 x 26 x 128		pad=1			
8 conv	256	3 x 3 / 1	26 x 26 x 128	-> 26 x 26 x 256		activation=leaky			
9 max	2	2 x 2 / 2	26 x 26 x 256	-> 13 x 13 x 256	[maxpool]	size=2			
10 conv	512	3 x 3 / 1	13 x 13 x 256	-> 13 x 13 x 512		stride=2			
11 max	2	2 x 2 / 1	13 x 13 x 512	-> 13 x 13 x 512	[convolutional]	batch_normalize=1			
12 conv	1024	3 x 3 / 1	13 x 13 x 512	-> 13 x 13 x 1024		filters=32			
13 conv	512	3 x 3 / 1	13 x 13 x 1024	-> 13 x 13 x 512		size=3			
14 conv	425	3 x 3 / 1	13 x 13 x 512	-> 13 x 13 x 425		stride=1			
15 detection						pad=1			

26

# Network configuration file

■ First node should be [net]

tiny-yolo-voc.cfg

layer	filters	size	input	output
0 conv	16	3 x 3 / 1	416 x 416 x 3	-> 416 x 416 x 16
1 max	2	2 x 2 / 2	416 x 416 x 16	-> 208 x 208 x 16
2 conv	32	3 x 3 / 1	208 x 208 x 16	-> 208 x 208 x 32
3 max	2	2 x 2 / 2	208 x 208 x 32	-> 104 x 104 x 32
4 conv	64	3 x 3 / 1	104 x 104 x 32	-> 104 x 104 x 64
5 max	2	2 x 2 / 2	104 x 104 x 64	-> 52 x 52 x 64
6 conv	128	3 x 3 / 1	52 x 52 x 64	-> 52 x 52 x 128
7 max	2	2 x 2 / 2	52 x 52 x 128	-> 26 x 26 x 128
8 conv	256	3 x 3 / 1	26 x 26 x 128	-> 26 x 26 x 256
9 max	2	2 x 2 / 2	26 x 26 x 256	-> 13 x 13 x 256
10 conv	512	3 x 3 / 1	13 x 13 x 256	-> 13 x 13 x 512
11 max	2	2 x 2 / 1	13 x 13 x 512	-> 13 x 13 x 512
12 conv	1024	3 x 3 / 1	13 x 13 x 512	-> 13 x 13 x 1024
13 conv	1024	3 x 3 / 1	13 x 13 x 1024	-> 13 x 13 x 1024
14 conv	125	1 x 1 / 1	13 x 13 x 1024	-> 13 x 13 x 125
15 detection				

```
[net]
batch=64
subdivisions=8
width=416
height=416
channels=3

learning_rate=0.001
max_batches =
40200
policy=steps
steps=-
1,100,20000,30000
scales=.1,10,.1,.1

[convolutional]
batch_normalize=1
filters=64
size=3
stride=1
pad=1
activation=leaky

[convolutional]
batch_normalize=1
filters=128
size=3
stride=1
pad=1
activation=leaky

[convolutional]
batch_normalize=1
filters=256
size=3
stride=1
pad=1
activation=leaky

[convolutional]
batch_normalize=1
filters=512
size=3
stride=1
pad=1
activation=leaky

[convolutional]
batch_normalize=1
filters=1024
size=3
stride=1
pad=1
activation=leaky

[convolutional]
batch_normalize=1
filters=1024
size=3
stride=1
pad=1
activation=leaky

[convolutional]
batch_normalize=1
filters=125
size=1
stride=1
pad=1
activation=linear

[region]
anchors =
1.08,1.19,
3.42,4.41,
6.63,11.38,
9.42,5.11,
16.62,10.52
bias_match=1
classes=20
coords=4
batch_normalize=1
num=5
softmax=1
jitter=.2
rescore=1
object_scale=5
noobject_scale=1
class_scale=1
coord_scale=1
absolute=1
thresh = .6
random=1
```

# YOLO

■ YOLO (You Only Look Once)

- ▶ An implementation of object-detection using Darknet.
- ▶ YOLO: <https://pjreddie.com/darknet/yolo/> → YOLO3
- ▶ YOLO (darknet): <https://pjreddie.com/darknet/yolov1/>
- ▶ YOLOv2 (darknet): <https://pjreddie.com/darknet/yolo/>
- ▶ YOLO (caffe): <https://github.com/xingwangsfu/caffe-yolo>
- ▶ YOLO (TensorFlow: Train+Test): <https://github.com/thtrieu/darkflow>
- ▶ YOLO (TensorFlow: Test): [https://github.com/gliese581gg/YOLO\\_tensorflow](https://github.com/gliese581gg/YOLO_tensorflow)



■ Why it is called YOLO

- ▶ A single neural network predicts bounding boxes and class probabilities directly from full images in one evaluation.

"You Only Look Once: Unified, Real-Time Object Detection", Joseph Redmon, Santosh Divvala, Ross Girshick, Ali Farhadi.



## YOLO on Caffe

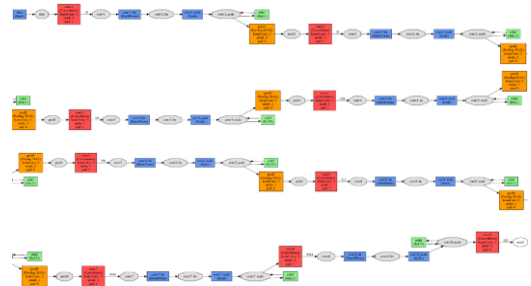
### Steps (still buggy)

- ▶ go to project directory
  - ➔ `$ cd work/codes/caffe_v1-projects/yolo_v2`
- ▶ converting Darknet configuration .cfg to Caffe .prototxt:
  - ➔ `$ python create_yolo_prototxt.py tiny-yolo.cfg yolo_tiny`
- ▶ convert Darknet weight .weights to Caffe model .caffemodel
  - ➔ `$ python create_yolo_caffemodel.py yolo_tiny_deploy.prototxt tiny-yolo.weights yolo_tiny.caffemodel`
- ▶ use training (pretrained) model
- ▶ running Yolo model
  - ➔ `$ python yolo_detect.py yolo_tiny_deploy.prototxt yolo_tiny.caffemodel images/dog.jpg`

refer to <https://github.com/tsingjinyun/caffe-yolov2>

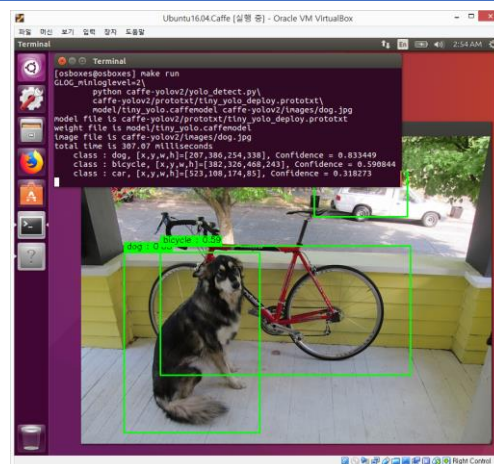
### Step in simple

- ▶ go to project directory
  - ➔ `$ cd work/codes/caffe_v1-projects/yolo_v2`
- ▶ get project
  - ➔ `$ make get`
- ▶ Run make
  - ➔ `$ make run`



31

## YOLO on Caffe



32



(주)퓨처디자인시스템

34051 대전광역시 유성구 문지로 193, KAIST 문지캠퍼스, F723호  
(042) 864-0211~0212 / [contact@future-ds.com](mailto:contact@future-ds.com) / [www.future-ds.com](http://www.future-ds.com)

Future Design Systems, Inc.

Faculty Wing F723, KAIST Munji Campus, 193 Munji-ro, Yuseong-gu, Daejeon 34051, Korea  
+82-042-864-0211~0212 / [contact@future-ds.com](mailto:contact@future-ds.com) / [www.future-ds.com](http://www.future-ds.com)



**FUTURE**  
Design Systems